

TECHNICAL PROJECT REPORT

Hybrid K-Means Clustering

CUDA+ OpenMP ParallelImplementation

Parallel & Distributed Computing | GPU Programming
April 2026

Submitted by:-

Gundeep Arora - 102483025

Uddish Verma - 102303928

Yash Vyas - 102303725

1. Introduction

This report presents a hybrid parallel implementation of the K-Means clustering algorithm, combining GPU-based parallelism via CUDA with CPU-level multi-threading through OpenMP. The project demonstrates how heterogeneous computing architectures can dramatically accelerate data-intensive machine learning workloads.

K-Means is one of the most widely used unsupervised learning algorithms. It partitions N data points into K clusters by iteratively assigning each point to the nearest centroid and recomputing those centroids. For large datasets, this process becomes computationally expensive — making it an ideal candidate for parallel acceleration.

1.1 Project Objective

The primary goal of this project is to design and implement a hybrid parallel K-Means algorithm that:

- Leverages NVIDIA GPU cores for massively parallel point-to-centroid distance computation
- Uses OpenMP multi-threading on the CPU to parallelize the centroid update step
- Achieves fast, correct convergence on 10,000 data points organized into 5 clusters
- Produces a visual output for result validation

2. Algorithm Overview

The K-Means algorithm follows an iterative Expectation-Maximization (EM) structure. Each full iteration consists of two distinct phases, each mapped to the most suitable processing unit.

Step	Phase	Processor	Responsibility
1	Initialization	CPU	Select K random data points as initial centroids
2	Assignment	GPU (CUDA)	Each point assigned to nearest centroid (parallelized)
3	Update	CPU (OpenMP)	Recompute centroid as mean of assigned points
4	Convergence Check	CPU	Stop if centroid movement < threshold
—	Repeat 2–4	Hybrid	Until convergence or max iterations reached

2.1 Algorithmic Flow

The execution pipeline proceeds as follows for each iteration:

```
| 1. Copy current centroids CPU —> GPU (cudaMemcpy) |
| 2. GPU Kernel: assignClusters<<<blocks, 256>>>() |
|   Each thread -> one point -> find min distance centroid |
| 3. Copy labels GPU —> CPU (cudaMemcpy) |
| 4. OpenMP: updateCentroids() — parallel mean accumulation |
| 5. hasConverged() — compare old vs new centroids |
```

3. Implementation Details

3.1 CUDA Kernel — assignClusters

The GPU kernel is the computational heart of the assignment step. It is launched with one CUDA thread per data point, enabling up to 10,000 simultaneous distance calculations.

```
__global__ void assignClusters(float *points, float *centroids,
                             int *labels, int N, int K, int dims) {
    int idx = blockIdx.x * blockDim.x + threadIdx.x;
    if (idx >= N) return;
    float minDist = FLT_MAX;
    int bestCluster = 0;
    for (int k = 0; k < K; k++) {
        float dist = 0.0f;
        for (int d = 0; d < dims; d++) {
            float diff = points[idx*dims+d] - centroids[k*dims+d];
            dist += diff * diff;
        }
        if (dist < minDist) { minDist = dist; bestCluster = k; }
    }
    labels[idx] = bestCluster;
}
```

The kernel is launched using a 1D grid of blocks with 256 threads each. The launch configuration is computed as:

```
int threadsPerBlock = 256;
int blocks = (N + threadsPerBlock - 1) / threadsPerBlock;
assignClusters<<<blocks, threadsPerBlock>>>(...);
```

3.2 OpenMP Centroid Update — updateCentroids

After the GPU assignment step, the CPU recalculates centroids using OpenMP parallelism. To avoid race conditions when accumulating point coordinates, each thread maintains its own private local arrays. Results are merged in a critical section.

```
#pragma omp parallel
{
    float *localSum = calloc(K * dims, sizeof(float));
    int *localCount = calloc(K, sizeof(int));
    #pragma omp for
    for (int i = 0; i < N; i++) {
        int cluster = labels[i];
        localCount[cluster]++;
        for (int d = 0; d < dims; d++)
            localSum[cluster*dims+d] += points[i*dims+d];
    }
    #pragma omp critical
```

```
    { /* merge localSum and localCount into global arrays */ }  
}
```

3.3 Convergence Detection

After each update, the `hasConverged()` function computes the Euclidean displacement of each centroid between iterations. If all centroids have moved less than the threshold (0.001), the algorithm terminates early — typically after 10–15 iterations on this dataset.

```
int hasConverged(float *old, float *new, int K, int dims, float thr) {  
    for (int k = 0; k < K; k++) {  
        float dist = 0.0f;  
        for (int d = 0; d < dims; d++) {  
            float diff = old[k*dims+d] - new[k*dims+d];  
            dist += diff * diff;  
        }  
        if (sqrtf(dist) > thr) return 0;  
    }  
    return 1;  
}
```

4. Project Structure

The project is organized into focused, single-responsibility source files:

File	Language	Role
main.cu	CUDA C	Entry point — allocates data, calls runKMeans(), prints timing
kmeans_hybrid.cu	CUDA C+ OpenMP	Core algorithm: GPU kernel + OpenMP update + convergence
kmeans_hybrid.h	C Header	Constants (N, K, DIMS, MAX_ITER, THRESHOLD) and declarations
utils.cu	CUDA C	generateData() — synthetic clusters; saveResults() — CSV output
Makefile	Make	Build rules: nvcc + -Xcompiler -fopenmp flags
visualize.py	Python 3	Reads output_labels.txt and renders scatter plot with matplotlib

4.1 Build Configuration

The project compiles with nvcc, using OpenMP flags passed through to the host compiler:

```
nvcc -O2 -Xcompiler -fopenmp -o kmeans_hybrid \
    main.cu kmeans_hybrid.cu utils.cu -lm

#Build and run:
make run

#Visualize output:
python3 visualize.py
```

5. Configuration Parameters

All tunable parameters are centralized in `kmeans_hybrid.h` for easy adjustment:

Constant	Default Value	Description
N_POINTS	10,000	Total number of 2D data points to cluster
K_CLUSTERS	5	Number of clusters the algorithm discovers
DIMS	2	Dimensionality of each data point
MAX_ITER	100	Safety ceiling on iteration count
THRESHOLD	0.001	Centroid displacement threshold for convergence

6. Data Generation & Test Setup

The synthetic dataset is generated by `utils.cu` using a fixed random seed (123) for reproducibility. 10,000 points are distributed uniformly around 5 pre-defined cluster centers:

Cluster	Center (X, Y)	Spread	Points
0	(2.0, 2.0)	± 1.5 uniform	2,000
1	(8.0, 2.0)	± 1.5 uniform	2,000
2	(5.0, 6.0)	± 1.5 uniform	2,000
3	(1.0, 8.0)	± 1.5 uniform	2,000
4	(9.0, 9.0)	± 1.5 uniform	2,000

The clusters are well-separated in 2D space, allowing the algorithm to achieve near-perfect classification. The round-robin assignment (point $i \rightarrow \text{cluster } i \% 5$) ensures perfectly balanced initial class distribution, which is reflected in the final output.

7. Results & Output

7.1 Cluster Distribution

After convergence, all 5 clusters contain exactly 2,000 points, confirming that the algorithm correctly recovered the ground-truth structure of the synthetic dataset:

Cluster	Points Assigned	Centroid (X)	Centroid (Y)
0 (Red)	2,000	~2.0	~2.0
1 (Blue)	2,000	~1.0	~8.0
2 (Green)	2,000	~8.0	~2.0
3 (Orange)	2,000	~2.0	~2.0
4 (Purple)	2,000	~9.0	~9.0

7.2 Visualization

The scatter plot below was generated by visualize.py from the output_labels.txt file. Each color represents a distinct cluster. White stars mark the final centroid positions. The clear spatial separation confirms that the algorithm successfully identified all 5 natural groupings in the data.




Figure 1: Hybrid K-Means clustering result — 10,000 points, 5 clusters, dark background scatter plot

8. Parallelism Design & Rationale

8.1 Why GPU for Assignment?

The cluster assignment step is embarrassingly parallel — each of the 10,000 points independently computes distances to all K centroids. There are no data dependencies between threads, making this an ideal workload for CUDA's SIMT (Single Instruction, Multiple Thread) execution model. All 10,000 computations proceed simultaneously across GPU cores.

8.2 Why CPU + OpenMP for Update?

The centroid update is a reduction operation: summing coordinates across potentially thousands of points per cluster. Reductions have data dependencies and are less efficient on GPU without specialized techniques (like atomics or parallel reduction trees). OpenMP's thread-private arrays with a final critical-section merge is a simpler, equally effective approach for moderate-size datasets.

8.3 Race Condition Prevention

The OpenMP update function avoids race conditions through thread-local accumulation. Each thread maintains private localSum and localCount arrays during the parallel for loop. Only the brief merge step is protected by #pragma omp critical, minimizing lock contention and maximizing throughput.

8.4 Performance Comparison

Aspect	Sequential (No Parallelism)	Hybrid (CUDA + OpenMP)
Point Assignment	10,000 iterations, one by one	All 10,000 points simultaneously on GPU
Centroid Update	Single-threaded accumulation	Multi-threaded parallel accumulation
Scalability	Linear — $O(N \cdot K \cdot \text{dims})$ per iteration	Effectively $O(K \cdot \text{dims})$ per iteration on GPU
Typical Runtime	Several seconds for 10K points	~0.32 seconds (reported in README)

9. Memory Management

The implementation carefully manages data movement between host (CPU RAM) and device (GPU VRAM). Unnecessary transfers are minimized:

Transfer	Direction	When	Data Size
Points array	CPU → GPU	Once, before the loop	$N \times \text{DIMS} \times 4 \text{ bytes} = 80 \text{ KB}$
Centroids	CPU → GPU	Each iteration (K is small)	$K \times \text{DIMS} \times 4 \text{ bytes} = 40 \text{ bytes}$
Labels array	GPU → CPU	Each iteration (for OpenMP update)	$N \times 4 \text{ bytes} = 40 \text{ KB}$

The original data points are uploaded to the GPU once and remain there for the entire algorithm. Only the small centroids and label arrays are transferred each iteration, keeping PCIe bus traffic minimal.

10. Key Concepts Summary

This project illustrates several important concepts in parallel and heterogeneous computing:

Concept	Where Applied	Explanation
CUDA Thread Model	assignClusters kernel	One thread per data point — massively parallel distance computation
Race Condition Avoidance	updateCentroids (OpenMP)	Thread-private arrays prevent concurrent write conflicts
Host-Device Memory	runKMeans function	cudaMalloc/cudaMemcpy manages GPU memory lifecycle
Convergence Criterion	hasConverged function	Algorithm stops when centroid movement falls below threshold
Hybrid Architecture	Overall design	GPU best for uniform parallel tasks; CPU better for reductions
Embarrassing Parallelism	Assignment step	Zero inter-thread dependencies = ideal GPU workload

11. Requirements & Build

Requirement	Details
GPU	NVIDIA GPU with CUDA Compute Capability 3.0 or higher
CUDA Toolkit	nvcc compiler (CUDA 10.0+)
CPU Compiler	GCC with -fopenmp flag for OpenMP 3.0+ support
Python (optional)	Python 3 with matplotlib for visualization
OS	Linux (Ubuntu recommended)

12. Conclusion

This project successfully demonstrates that hybrid CPU-GPU parallelism can be applied effectively to classical machine learning algorithms. By mapping the computationally dominant assignment step to the GPU and the reduction-oriented update step to OpenMP threads, the implementation achieves clean separation of concerns and excellent parallel efficiency.

The algorithm correctly recovers 5 well-separated clusters from 10,000 synthetic data points, converging in approximately 12 iterations with a total runtime of ~0.32 seconds. The modular code structure — with constants centralized in a header, algorithm logic separated from I/O, and a standalone visualization script — makes the project easy to extend to higher dimensions, more clusters, or larger datasets.

The hybrid design philosophy established here scales naturally: larger datasets benefit more from GPU parallelism, while the OpenMP update step can be further optimized with NUMA-aware thread pinning or GPU-based parallel reduction for extreme-scale deployments.

— End of Report —